

Too many knots

Abstract

Introduction

Recent computational advances have allowed for the rapid estimation of large spatial and spatiotemporal models in ecology and related fields (Banerjee & Fuentes 2012). Motivating questions for these applications to large georeferenced datasets including tracking abundance in time and space, identifying environmental predictors of occurrence, or quantifying range shifts in distribution (Ross *et al.* 2012; Miller *et al.* 2013; Thorson & Barnett 2017).

There are a number of statistical approaches for developing models with spatial components; examples include generalized additive models (GAMs) (Murase *et al.* 2009; Forney *et al.* 2012), non-parametric approaches (Mutanga *et al.* 2012), and applications of Gaussian random fields (Latimer *et al.* 2009; Datta *et al.* 2016). Advances by (Rue *et al.* 2009) have enabled the development of the Integrated Nested Laplace Approximation (INLA) as an fast and efficient approximation to estimating high-dimensional Gaussian random fields with INLA. These models have been extended to include the ability to fit geographic barriers, such as islands or lakes, anisotropic covariance functions, and non-Gaussian data (Bakka *et al.* 2018). In addition to estimating Bayesian posterior distributions (Lindgren & Rue 2015), INLA has been integrated with Template Model Builder (TMB; (Kristensen *et al.* 2016; Osgood-Zimmerman & Wakefield 2021)) to enable auto-differentiation – the latter being adopted in many spatial fisheries applications around the world (Thorson 2019b; Johnson *et al.* 2019).

Constructing spatiotemporal models with INLA or related software involves many of the same considerations as with simpler non-spatial model, e.g. “Which covariates are most important?” or “Does my model do a satisfactory job of making predictions?” Spatial or spatiotemporal models may include additional considerations, such as specifying the form of spatiotemporal variation (Anderson & Ward 2019). One of the more complicated and important decisions – particularly for INLA based models – lies in specifying the dimensionality of the latent spatial fields (Lindgren & Rue 2015). Unlike GAMs, which penalize smooth terms (cite wood 2003?), penalties in overfitting INLA models are not as apparent.

One of the first steps in fitting a model with INLA involves choosing the appropriate resolution of the spatial fields, using meshes (maps consisting of small triangles, e.g. Figure 1). By estimating values of the spatial field at the triangle vertices (or ‘knots’), the spatial field may be triangulated to coordinates where data are observed (Lindgren *et al.* 2011). There are a number of decision points to be made in constructing the mesh, including whether to include a boundary, whether the boundary is regular, the maximum edge of any triangle (in the core or boundary area), and the minimum angles allowed by the triangles (Gómez-Rubio). Each of these decisions can impact the density and number of small triangles that make up the mesh – and in turn, the accuracy of predictions. Though the mesh construction is important, this step has been omitted from best practices guides using the INLA based approach (Thorson 2019a), and we’re only aware of one sensitivity analysis examining the importance of mesh construction (Righetto *et al.* 2020). While increasing the resolution of the mesh has been suggested to improve accuracy (Godana *et al.* 2019), other authors have also noted that finer meshes don’t always translate into better predictions (Huang *et al.* 2017).

The objective of this paper is to demonstrate that higher resolution spatial approximations in INLA may not translate to better predictions. Using a series of sensitivity analyses on a real world case study of groundfish data in the Northeast Pacific Ocean, we evaluate the out of sample predictive performance as a function of mesh resolution (number of vertices). We further evaluate these results across a range of cross validation sampling schemes, including an extreme case of leave one out cross validation (LOOCV). All data for our analysis are publicly available, and code and data to reproduce our analysis is on Github (XXXX).

Methods

Data

As an case study to evaluate the effect of mesh resolution on predictive ability of INLA based models, we use a publicly available dataset of catch rates from scientific surveys in the Northeast Pacific Ocean (<https://www.nwfsc.noaa.gov/data>). The West Coast Bottom Trawl Survey (WCBTS) dataset represents a fishery-independent survey collected annually over a 16-year period, 2003-2018, covering the California Current ecosystem on the west coast of the USA (Keller *et al.* 2017). The survey is designed to estimate the abundance, sizes, and age composition of commercially and recreationally targeted groundfish species found in near-bottom habitats. The fundamental methods of the survey, sampling intensity, gear, seasonal and spatial coverage, and random stratified design have remained relatively constant. Because of the large number of model runs needed, we restricted our analysis to only use data from 2018 ($n = 702$ observations). Though hundreds of species our encountered in the survey, we also restrict our analysis to only include data from four species that are of interest to fisheries, but also represent a range of occurrences: dover

sole (*Microstomus pacificus*, encountered in 593 tows), petrale sole (*Eopsetta jordani*, 291 tows), lingcod (*Ophiodon elongatus*, 225 tows), and darkblotched rockfish (*Sebastes crameri*, 105 tows).

Modeling

Because of zero-inflation in the fish survey datasets like the WCBTS data, conventional approaches to analyzing this data have used a delta or hurdle model to separately model occurrence and positive catches (Thorson 2019a; Johnson *et al.* 2019). We adopt a similar approach here, using a Binomial distribution (logit link) to model species occurrences, and a Gamma distribution (log link) to model positive catch rates. For simplicity and consistency with fisheries assessments using these data, we fit models estimating an intercept and spatial field throughout, e.g.

$$g(u_s) = B_0 + \omega_s$$

where ω_s is the estimated spatial field.

All of the meshes in our analysis used a non-convex boundary and offset based on 5% of the data (Figure 1). We fixed the maximum edge of triangles in the mesh at 100km in the core area and 500km closer to the boundary (Figure 1). Throughout our analyses, we adjust the mesh resolution by varying the cutoff distance (representing the distance where 2 points are treated as a single vertex) because that has been shown to have the greatest impact on mesh construction (Righetto *et al.* 2020). All models were fit using R-INLA (Lindgren & Rue 2015; R Core Team 2021) and specifically the ‘inlabru’ package, which is a convenient wrapper for the INLA software (Bachl *et al.* 2019). Penalized complexity priors were used on the Matérn covariance parameters (Simpson *et al.* 2015; Osgood-Zimmerman & Wakefield 2021). Code to replicate our analysis is provided on Github [link XXXX].

Sensitivities

As a preliminary evaluation of whether the spatial resolution of INLA based models can result in overfitting, we evaluated 40 cutoff values for each species and hurdle sub-model included in our analysis (cutoff values ranging from 3-120km). There is an inverse relationship between cutoff values and mesh resolution or vertices; while the results varied slightly by species, our meshes ranged in size from 75 - 1348 vertices. The high resolution meshes with more vertices than observations were included as an upper extreme, though many applications of INLA to similar survey data have used 500-600 vertices (Thorson 2019a; Tolimieri *et*

86 *al.* 2020).

87 To evaluate model performance, we conducted 10-fold cross-validation. Though not a 1-dimensional dataset,
88 the strong latitude gradient has driven spatial stratification in previous analyses (Keller *et al.* 2017) (Figure
89 1). We a priori divided the survey domain into 10 spatial blocks by latitude quantiles (with approximately
90 the same number of points in each box). Each model was run using *inlabru* (@ Bachl *et al.* 2019), and used
91 to predict the data for the held out test data. The predictive density was calculated using the binomial or
92 Gamma densities, and total predictive densities were calculated by summing across folds.

93 As a second sensitivity, we evaluated whether the particular cross-validation method used may be influencing
94 the results. Recent work has demonstrated that for spatially structured data, creating random folds that
95 ignore spatial autocorrelation may bias results (Valavi *et al.* 2019). Using the presence-absence data from
96 Dover sole, we examined the cross validation performance across a range of mesh resolutions (cutoff values
97 ranging from 15-75, or 122 - 594 vertices), number of spatial blocks (19 to 218) and number of folds used for
98 cross validation (10, 40). Selection of spatial blocks was done using the ‘blockCV’ package in R (Valavi *et*
99 *al.* 2019) with varying distance bands between test and training sets (25 - 150km).

100 To further validate the inference from the spatial cross-validation, we examined the performance of out of
101 sample approximations in INLA. We focused on the conditional predictive ordinate (CPO), which is an
102 approximation to the leave-one-out cross-validation score (LOOCV). In an ideal world, we would use cross
103 validation to evaluate the out of sample predictive accuracy of a model. The total predictive log density can
104 be represented as the sum of individual log posterior probabilities, $\sum_{i=1}^n \log[P(Y_i|Y_{\neq i})]$, where $P(Y_i|Y_{\neq i})$ is
105 the probability of test data i . Because conducting true leave one out cross-validation may be prohibitive for
106 large datasets, INLA approximates the out of sample posterior probability for each observation. Previous
107 research has shown that the CPO approximation performs similarly to MCMC cross-validation (Held *et*
108 *al.* 2010), yet it is less clear whether CPO is similar to true cross validation in comparing mesh resolution.
109 Focusing on the presence-absence model of Dover sole (Figure 2), we first calculated the total CPO score
110 across a range of cutoff values. For comparison, we then calculated the true predictive density using LOOCV
111 across a similar range of cutoffs (fitting the INLA model 702 times for each value).

112 Results

113 Fig. 1 just demonstrates the data used in the study, and a couple mesh configurations (coarse and fine,
114 cutoff of 20 and 80, translating to 394 and 110 vertices)

115 Fig. 2 shows the decline in performance as a function of large knots, across the 4 species. Similarly, Fig S1

shows this for the gamma distribution.

Fig. 3 shows the decline in performance as a function of large knots happens regardless of the type of sampling used (sensitivity to # blocks and # folds).

Fig. 4. shows the comparing between CPO (via INLA) and real LOOCV. LOOCV shows that performance declines as knots become too large, but CPO keeps increasing

Fig. 5 Illustrates downstream effects of knots on estimates of downstream parameters ()

Discussion


Larger meshes aren't necessarily better. Others have pointed this out, https://punama.github.io/BDI_INLA/, but it gets overlooked


Sub-optimal meshes are likely to have an impact on derived things we care about, like indices of abundance – both bias and precision, fig s2

Datasets like the WCTBS support an optimal mesh of ~ 300 or fewer knots

We recommend analysts evaluate models across a range of resolutions and carefully consider what cross-validation scheme may be most appropriate for the question at hand

Figures






134

135

136



137

138

139




Figure 1: Figure 4




Figure 2: Figure S1




Figure 3: Figure S2

References

- Anderson, S.C. & Ward, E.J. (2019). Black swans in space: Modeling spatiotemporal processes with extremes. *Ecology*, **100**, e02403.
- Bachl, F.E., Lindgren, F., Borchers, D.L. & Illian, J.B. (2019). Inlabru: An R package for Bayesian spatial modelling from ecological survey data. *Methods in Ecology and Evolution*, **10**, 760–766.
- Bakka, H., Rue, H., Fuglstad, G.-A., Riebler, A., Bolin, D., Illian, J., Krainski, E., Simpson, D. & Lindgren, F. (2018). Spatial modeling with R-INLA: A review. *WIREs Computational Statistics*, **10**, e1443.
- Banerjee, S. & Fuentes, M. (2012). Bayesian modeling for large spatial datasets. *WIREs Computational Statistics*, **4**, 59–66.
- Datta, A., Banerjee, S., Finley, A.O. & Gelfand, A.E. (2016). Hierarchical Nearest-Neighbor Gaussian Process Models for Large Geostatistical Datasets. *Journal of the American Statistical Association*, **111**, 800–812.
- Forney, K.A., Ferguson, M.C., Becker, E.A., Fiedler, P.C., Redfern, J.V., Barlow, J., Vilchis, I.L. & Ballance, L.T. (2012). Habitat-based spatial models of cetacean density in the eastern Pacific Ocean. *Endangered Species Research*, **16**, 113–133.
- Godana, A.A., Mwalili, S.M. & Orwa, G.O. (2019). Dynamic spatiotemporal modeling of the infected rate of visceral leishmaniasis in human in an endemic area of Amhara regional state, Ethiopia. *PLOS ONE*, **14**, e0212934.
- Gómez-Rubio, V. *Bayesian inference with INLA*.
- Held, L., Schrödle, B. & Rue, H. (2010). Posterior and Cross-validatory Predictive Checks: A Comparison of MCMC and INLA. *Statistical Modelling and Regression Structures: Festschrift in Honour of Ludwig Fahrmeir* (eds T. Kneib & G. Tutz), pp. 91–110. Physica-Verlag HD, Heidelberg.
- Huang, J., Malone, B.P., Minasny, B., McBratney, A.B. & Triantafyllis, J. (2017). Evaluating a Bayesian modelling approach (INLA-SPDE) for environmental mapping. *The Science of the Total Environment*, **609**, 621–632.
- Johnson, K.F., Thorson, J.T. & Punt, A.E. (2019). Investigating the value of including depth during spatiotemporal index standardization. *Fisheries Research*, **216**, 126–137.
- Keller, A.A., Wallace, J.R. & Methot, R.D. (2017). *The Northwest Fisheries Science Center’s West Coast Groundfish Bottom Trawl Survey: History, Design, and Description*. Seattle, WA.

- Kristensen, K., Nielsen, A., Berg, C.W., Skaug, H. & Bell, B.M. (2016). TMB: Automatic Differentiation and Laplace Approximation. *Journal of Statistical Software*, **70**, 1–21.
- Latimer, A.M., Banerjee, S., Jr, H.S., Mosher, E.S. & Jr, J.A.S. (2009). Hierarchical models facilitate spatial analysis of large data sets: A case study on invasive plant species in the northeastern United States. *Ecology Letters*, **12**, 144–154.
- Lindgren, F. & Rue, H. (2015). Bayesian Spatial Modelling with R-INLA. *Journal of Statistical Software*, **63**, 1–25.
- Lindgren, F., Rue, H. & Lindström, J. (2011). An explicit link between Gaussian fields and Gaussian Markov random fields: The stochastic partial differential equation approach. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, **73**, 423–498.
- Miller, D.L., Burt, M.L., Rexstad, E.A. & Thomas, L. (2013). Spatial models for distance sampling data: Recent developments and future directions. *Methods in Ecology and Evolution*, **4**, 1001–1010.
- Murase, H., Nagashima, H., Yonezaki, S., Matsukura, R. & Kitakado, T. (2009). Application of a generalized additive model (GAM) to reveal relationships between environmental factors and distributions of pelagic fish and krill: A case study in Sendai Bay, Japan. *ICES Journal of Marine Science*, **66**, 1417–1424.
- Mutanga, O., Adam, E. & Cho, M.A. (2012). High density biomass estimation for wetland vegetation using WorldView-2 imagery and random forest regression algorithm. *International Journal of Applied Earth Observation and Geoinformation*, **18**, 399–406.
- Osgood-Zimmerman, A. & Wakefield, J. (2021). A Statistical Introduction to Template Model Builder: A Flexible Tool for Spatial Modeling.
- R Core Team. (2021). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria.
- Righetto, A.J., Faes, C., Vandendijck, Y. & Ribeiro, P.J. (2020). On the choice of the mesh for the analysis of geostatistical data using R-INLA. *Communications in Statistics - Theory and Methods*, **49**, 203–220.
- Ross, B.E., Hooten, M.B. & Koons, D.N. (2012). An Accessible Method for Implementing Hierarchical Models with Spatio-Temporal Abundance Data. *PLOS ONE*, **7**, e49395.
- Rue, H., Martino, S. & Chopin, N. (2009). Approximate Bayesian inference for latent Gaussian models by using integrated nested Laplace approximations. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, **71**, 319–392.

- 198 Simpson, D.P., Rue, H., Martins, T.G., Riebler, A. & Sørbye, S.H. (2015). Penalising model component
199 complexity: A principled, practical approach to constructing priors. *arXiv:1403.4630 [stat]*.
- 200 Thorson, J.T. (2019a). Guidance for decisions using the Vector Autoregressive Spatio-Temporal (VAST)
201 package in stock, ecosystem, habitat and climate assessments. *Fisheries Research*, **210**, 143–161.
- 202 Thorson, J.T. (2019b). Measuring the impact of oceanographic indices on species distribution shifts: The
203 spatially varying effect of cold-pool extent in the eastern Bering Sea. *Limnology and Oceanography*, **64**,
204 2632–2645.
- 205 Thorson, J.T. & Barnett, L.A.K. (2017). Comparing estimates of abundance trends and distribution shifts
206 using single- and multispecies models of fishes and biogenic habitat. *ICES Journal of Marine Science*,
207 **74**, 1311–1321.
- 208 Tolimieri, N., Wallace, J. & Haltuch, M. (2020). Spatio-temporal patterns in juvenile habitat for 13 ground-
209 fishes in the California Current Ecosystem. *PLOS ONE*, **15**, e0237996.
- 210 Valavi, R., Elith, J., Lahoz-Monfort, J.J. & Guillera-Arroita, G. (2019). blockCV: An r package for generat-
211 ing spatially or environmentally separated folds for k-fold cross-validation of species distribution models.
212 *Methods in Ecology and Evolution*, **10**, 225–232.